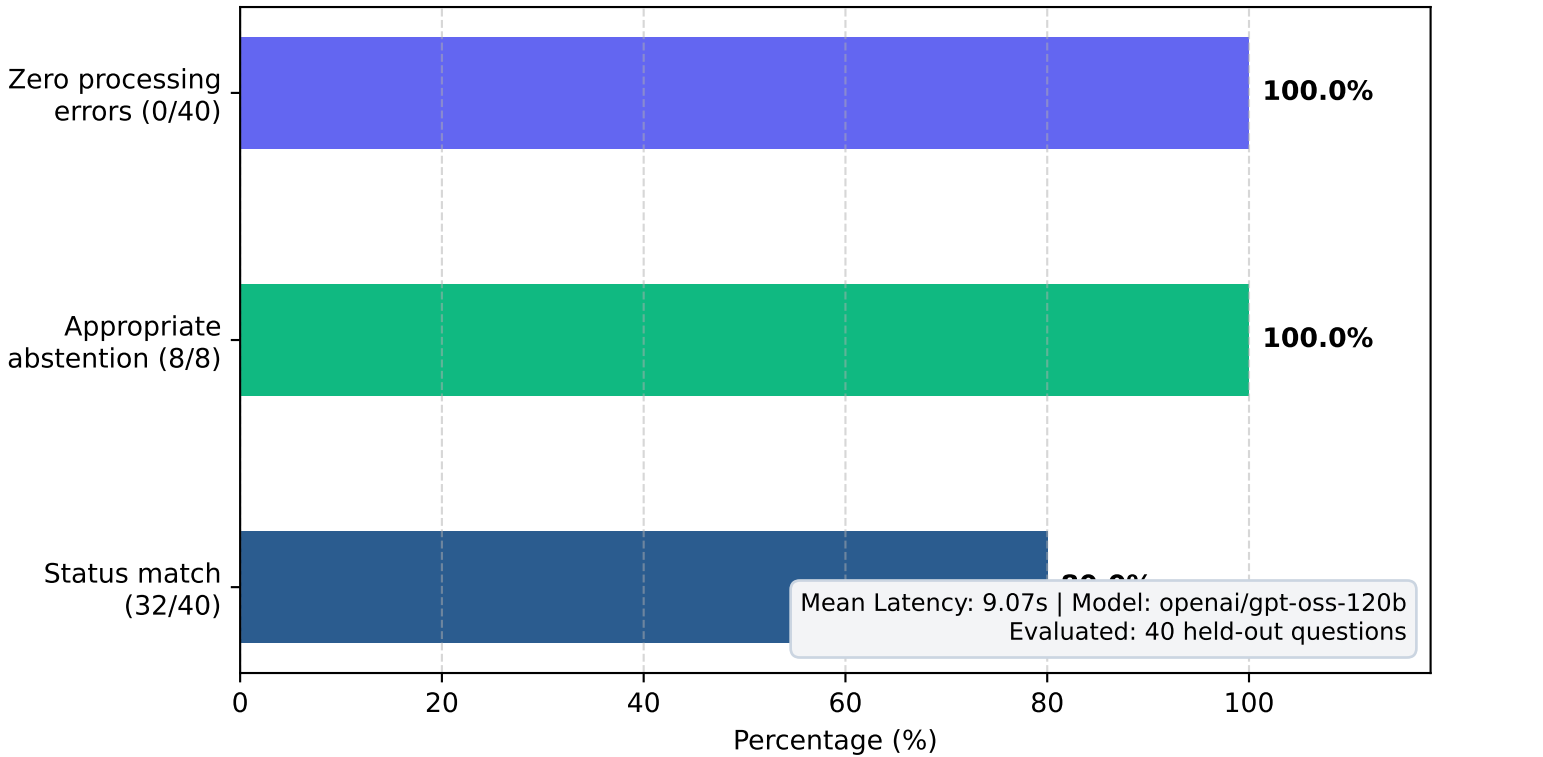
Held-Out Benchmark: Variant A Preliminary Result
(Preliminary generalisation check only; Variants B, C, and D were not evaluated on the held-out set)
Zero processing errors (0/40)
100.0%
Appropriate abstention (8/8)
100.0%
Status match (32/40)
Mean Latency: 9.07s | Model: openai/gpt-oss-120b
Evaluated: 40 held-out questions
0
20
40
60
80
100
Percentage (%)