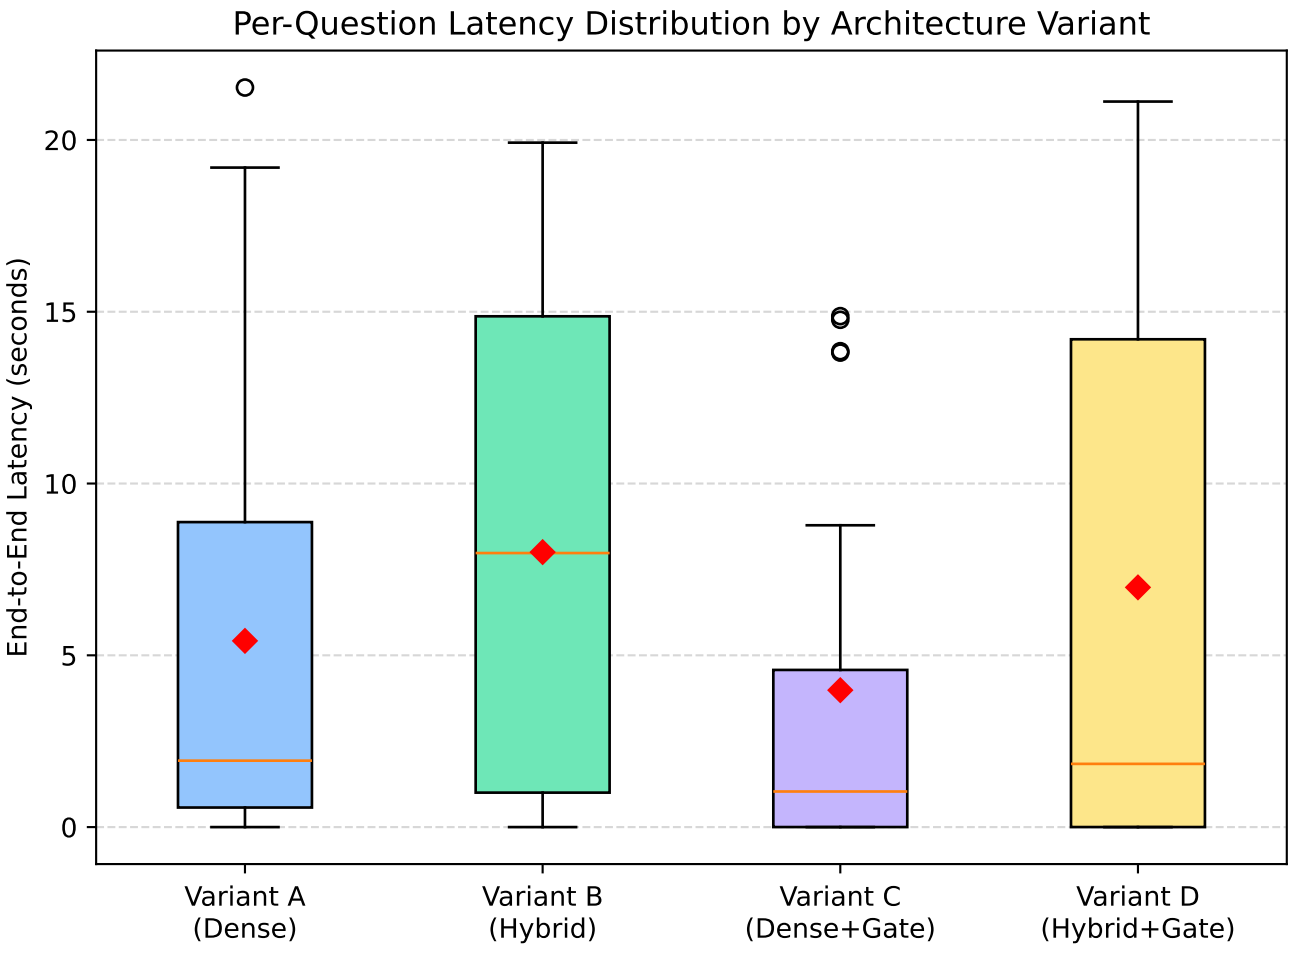
Per-Question Latency Distribution by Architecture Variant
End-to-End Latency (seconds)
0
5
10
15
20
Variant A
(Dense)
Variant B
(Hybrid)
Variant C
(Dense+Gate)
Variant D
(Hybrid+Gate)